

Report

Table of Contents

1. Purpose
2. Two main approaches
3. Data Cleaning and Merging Pipeline
4. Model Evaluation Criteria
5. Models Tried
 - RandomForestClassifier on flat dataset
 - Target Variable
 - Confusion Matrix
 - Classification Report
 - RandomForestClassifier on panel dataset
 - LGBMClassifier
 - Target Variable
 - Model Results (All Months)
 - Model Results (Months 1–12 only)
 - Monthly Fire Risk Prediction Using XGBoost
 - Script Location
 - Data Pipeline
 - Feature Engineering
 - Model
 - Evaluation
 - Threshold Optimization
 - Final Model Evaluation
 - Recommendations
 - Binary vs Probabilistic Forecasting
 - Summary

Purpose

Project A – Predicting High Fire Risk Areas in Montreal

Objective: The objective of this project is to predict high fire risk areas by month in the city of Montreal, based on historical firefighter intervention data and additional open datasets.

Two main approaches

“Panel” Approach

Structure: Each building has one row per month (e.g., Jan 2018, Feb 2018, ..., Dec 2022).

ID_UEV	YEAR	MONTH	Features...	fire_occurred
001	2020	1	...	0

```

001    | 2020 |    2    | ...    | 0
001    | 2020 |    3    | ...    | 1
...

```

`fire_occurred = 1` for the month of fire; 0 otherwise.

No `fire_month` column needed — you’re modeling fire per building-month.

Good for temporal modeling, like survival analysis or monthly fire risk.

Fire Month or 13 Per-Building Row Approach

Structure: One row per building; `fire_month = 1–12` if fire occurred, else 13.

```

ID_UEV | Features... | fire_month
-----|-----|-----
001    | ...          | 3
002    | ...          | 13
003    | ...          | 7
...

```

Simple classification task: which month is the fire, or no fire (13).

Only one row per building — no temporal unfolding.

`fire_month = 13` explicitly encodes “no fire”.

Data Cleaning and Merging Pipeline

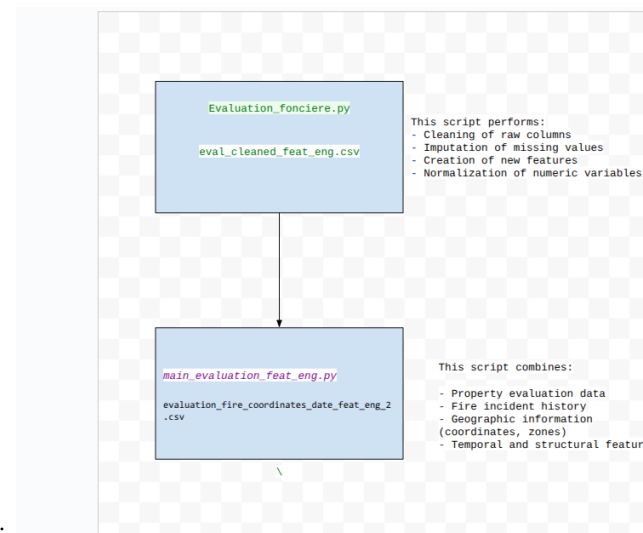


Diagram showing the data cleaning and merging pipeline:

Model evaluation criteria - Priority is given to recall over precision as we’d rather capture more fire risk including a few false negatives than miss high risk buildings - Train set/Test set : we used a temporal split rather than random split Train = data before 2024 Test = data of 2024 (full year)

This is especially important because of the approximation of fire location due to data obfuscation to avoid leaking knowledge of fires into the train set.

Models tried

RandomForestClassifier on flat dataset

(Located in file EDA-incident-evaluation-fonciere.ipynb, for pipeline see instructions)

Target Variable

$$Y = P(\text{Fire} \mid X)$$

Where X includes:

- log_terrain
- log_batiment
- log_etage_hors_sol
- log_nombre_de_logement
- ANNEE_CONSTRUCTION
- density

Confusion Matrix

	Predicted False	Predicted True
Actual False	64,073	9,540
Actual True	8,273	50,871

Classification Report

Label	Precision	Recall	F1-score	Support
False	0.89	0.87	0.88	73,613
True	0.84	0.86	0.85	59,144
Accuracy			0.87	132,757
Macro avg	0.86	0.87	0.86	132,757
Weighted avg	0.87	0.87	0.87	132,757

ROC AUC: 0.936

Target variable

$$Y = P(\text{Fire month} \mid X)$$

Confusion Matrix

- NOMBRE_LOGEMENT
- ANNEE_CONSTRUCTION
- SUPERFICIE_TERRAIN
- SUPERFICIE_BATIMENT
- LONGITUDE
- LATITUDE

Note: Properties with no recorded fire were assigned to month 13 to indicate the absence of fire incidents

Metric	Precision	Recall	F1-score	Support
Accuracy			0.676	132,757
Macro avg	0.457	0.348	0.388	132,757
Weighted avg	0.629	0.676	0.640	132,757

I then tried to train only on months 1-12:

Metric	Precision	Recall	F1-score	Support
Accuracy			0.434	58,954
Macro avg	0.439	0.433	0.434	58,954
Weighted avg	0.437	0.434	0.433	58,954

Monthly Fire Risk Prediction Using XGBoost

Script Location

/dataprep/time_model_Xgboost.ipynb

This script trains and evaluates a binary classifier to predict whether a fire will occur in a specific building in a given month. It utilizes a **dense panel dataset** with rich building-level and temporal-spatial features.

Data Pipeline

- **Input file:** building_month_fire_panel_feat_eng.csv
 - **Granularity:** Monthly panel of all buildings
 - **Target Variable:** HAS_FIRE_THIS_MONTH (0 or 1)
-

Feature Engineering

Structural & Geographic Features:

- MUNICIPALITE, ETAGE_HORS_SOL, NOMBRE_LOGEMENT, AGE_BATIMENT
- SUPERFICIE_TERRAIN, SUPERFICIE_BATIMENT, RATIO_SURFACE, DENSITE_LOGEMENT
- HAS_MULTIPLE_LOGEMENTS, CODE_UTILISATION, CATEGORIE_UEF
- NO_ARROND_ILE_CUM, BUILDING_COUNT

Zone-Level Fire Risk:

- FIRE_FREQUENCY_ZONE, FIRE_RATE_ZONE, FIRE_COUNT_LAST_YEAR_ZONE
- FIRE_RATE_ZONE_NORM, FIRE_COUNT_LAST_YEAR_ZONE_NORM

Temporal Lag Features:

- fire_last_1m, fire_last_2m, fire_last_3m
- fire_cumcount, fire_rolling_3m, fire_rolling_6m, fire_rolling_12m
- month_num, year

Model: XGBoostClassifier

- Handles class imbalance with `scale_pos_weight`
- Supports categorical variables with `enable_categorical=True`
- Optimized with:
 - `n_estimators=200`, `max_depth=6`, `learning_rate=0.1`
 - `subsample=0.8`, `colsample_bytree=0.8`

Evaluation (Default Threshold = 0.5)

Class	Precision	Recall	F1-score	Support
No Fire (0)	0.9903	0.7436	0.8494	3,674,405
Fire (1)	0.0243	0.4665	0.0461	50,239

- **Accuracy:** 73.99%
- **Macro F1:** 0.4477
- **Weighted F1:** 0.8379

Precision for fires is very low, but recall is moderate. Useful for prioritization, not alarms.

Threshold Optimization

Evaluated thresholds: $0.2 \rightarrow 0.55$ - Best **F2 Score** (recall-focused): **0.55**

Threshold	Precision	Recall	F2 Score
0.50	0.027	0.603	0.113
0.55	0.0262	0.378	0.103

Final Model Evaluation @ Threshold = 0.55

Class	Precision	Recall	F1-score	Support
No Fire	0.9896	0.8081	0.8897	3,674,405
Fire	0.0262	0.3780	0.0490	50,239

- **Confusion Matrix:**
 - True Negatives: ~2.97M
 - False Positives: ~700k
 - False Negatives: ~31k
 - True Positives: ~19k

High false positives but acceptable for early warning.

Recommendations

- Use **threshold = 0.50–0.55** depending on recall vs precision preference
 - **Precision is low**, so interpret predictions as risk levels
 - Use **predicted probabilities** for prioritizing inspections
 - Add weather, crime, or inspection features for improved separation
-

Binary vs Probabilistic Forecasting

Aspect	<code>model.predict()</code> (Binary)	<code>model.predict_proba()</code> + Threshold
Nature	Hard decision (0 or 1)	Probabilistic risk score (0.0–1.0)
Control over recall/precision	Fixed at 0.5	Fully tunable
Risk prioritization	Not possible	Buildings can be ranked

Aspect	<code>model.predict()</code> (Binary)	<code>model.predict_proba()</code> + Threshold
Early warning usability	Rigid	Flexible, interpretable
Best suited metric	Accuracy, F1	F2 Score, Precision@K, Recall

Recommendation: Use predicted probabilities to forecast fire risk, prioritize inspections, and adapt to seasonal or strategic needs.

Next Steps

- Deploy as a **ranking tool** not a strict classifier
- Create dashboards that visualize monthly fire risk scores by building
- Use **precision@K** and **F2-score** as main evaluation metrics

Summary

XGBoost + panel-level fire features + lag history yields a decent early-warning fire detection model. Prioritize **recall and interpretability**, refine with more context over time. ““”